

Wei Lin (Mr.)

Updated October 2026

Email: wlin2021at@gmail.com

HomePage: [link](#)

GitHub: [link](#)

Google Scholar: [link](#)

Research interests

Reliable Agent Systems, Multimodal LLMs, Domain Adaptation, Video Understanding

Work Experience

Staff Applied Scientist
Xero UK Ltd
London, UK
Jan 2026 – Present

Postdoc at the Institute for Machine Learning
(headed by Prof. Sepp Hochreiter, Father of LSTM)
Johannes Kepler University
Linz, Austria
Oct 2023 – Dec 2025

Research Assistant
Institute of Computer Graphics and Vision
Graz University of Technology
Graz, Austria
Jan 2019 – Sep 2023

Education

Graz University of Technology
PhD (Dr.techn) in Computer Science
Mentors:
Prof. Horst Bischof (Graz University of Technology)
Prof. Hilde Kuehne (University of Tuebingen, MIT-IBM Watson AI lab)
Graz, Austria
Jan 2019 – Mar 2024

Technical University of Munich
M.Sc. in Electrical and Computer Engineering
Mentor: Prof. Eckehard Steinbach
Munich, Germany
Oct 2015 - Dec 2018

Organization

Program Chair of the 2nd, 3rd and 4th Workshop on "What is Next in Multimodal Foundation Models" on CVPR 2024, 2025 and ICCV 2025, **Challenge Chair** for the MMFM-Challenge of Multimodal Models on Document Understanding on CVPR 2024

Publications

Architectural Sampling: Test-Time Scaling via Computational Diversity in Frozen Vision-Language Models
Akshit Singh, Shyam Marjit, Wei Lin, Leonid Karlinsky, M. Jehanzeb Mirza
[arXiv 2026](#)

Visual Grounding First, Multimodal In-context Learning Follows

*Minhyuk Seo, *Minjae Lee, Chaeun Lee, Wei Lin, Muhammad Jehanzeb Mirza, Tinne Tuytelaars, Jonghyun Choi (*equal contribution)
NeurIPS 2026

Efficient Encoder-Free Fourier-based 3D Large Multimodal Model

Guofeng Mei, Wei Lin, Luigi Riz, Yujiao Wu, Yiming Wang, Fabio Poiesi
CVPR 2026

TTRV: Test-Time Reinforcement Learning for Vision Language Models

Akshit Singh, Shyam Marjit, Wei Lin, Paul Gavrikov, Serena Yeung-Levy, Hilde Kuehne, Rogerio Feris, Sivan Doveh, James Glass, M. Jehanzeb Mirza
CVPR 2026

VisualOverload: Probing Visual Understanding of VLMs in Really Dense Scenes

Paul Gavrikov, Wei Lin, M. Jehanzeb Mirza, Soumya Jahagirdar, Muhammad Huzaifa, Sivan Doveh, Serena Yeung-Levy, James Glass, Hilde Kuehne
CVPR 2026

PRISMM-Bench: A Benchmark of Peer-Review Grounded Multimodal Inconsistencies

Lukas Selch, Yufang Hou, M. Jehanzeb Mirza, Sivan Doveh, James Glass, Rogerio Feris, Wei Lin
ICLR 2026

AViON4D: Audio-Visual Open-Vocabulary 4D Egocentric Scene Understanding

Irene Ballester, Pedro Hermosilla, Wei Lin, James R. Glass, M. Jehanzeb Mirza, Martin Kampel
CVPR 2026 OpenSUN3D Workshop

Probing the Effectiveness of World Models for Spatial Reasoning through Test-time Scaling

Saurav Jha, M. Jehanzeb Mirza, Wei Lin, Shiqi Yang, Sarath Chandar
World Modeling Workshop 2026

The RoboSense Challenge: Sense Anything, Navigate Anywhere, Adapt Across Platforms

Lingdong Kong, Shaoyuan Xie, Zeying Gong, Ye Li, Meng Chu, Ao Liang, Yuhao Dong, et al. (including Wei Lin)
Official IROS 2025 RoboSense Challenge Report

From Generated Human Videos to Physically Plausible Robot Trajectories

James Ni, Zekai Wang, Wei Lin, Amir Bar, Yann LeCun, Trevor Darrell, Jitendra Malik, Roei Herzig

arXiv 2025

GBlobs: Local LiDAR Geometry for Improved Sensor Placement Generalization

Dušan Malić, Christian Fruhwirth-Reisinger, Alexander Prutsch, Wei Lin, Samuel Schulter, Horst Possegger

1st place, IROS 2025 RoboSense Challenge, Track #3: Cross-Sensor Placement 3D Object Detection

Exploring Modality Guidance to Enhance VFM-based Feature Fusion for UDA in 3D Semantic Segmentation

Johannes Spoecklberger, Wei Lin, Pedro Hermosilla, Sivan Doveh, Horst Possegger, M. Jehanzeb Mirza

CVPR 2025 DriveX Workshop

pLSTM: parallelizable Linear Source Transition Mark networks

Korbinian Pöppel, Richard Freinschlag, Thomas Schmied, Wei Lin, Sepp Hochreiter

NeurIPS 2025

STSBench: A Spatio-temporal Scenario Benchmark for Multi-modal Large Language Models in Autonomous Driving

Christian Fruhwirth-Reisinger, Dušan Malić, Wei Lin, David Schinagl, Samuel Schulter, Horst Possegger

NeurIPS 2025 Datasets & Benchmarks Track

Teaching VLMs to Localize Specific Objects from In-context Examples

Sivan Doveh, Nimrod Shabtay, Wei Lin, Eli Schwartz, Hilde Kuehne, Raja Giryes, Rogerio Feris, Leonid Karlinsky, James Glass, Assaf Arbelle, Shimon Ullman, Muhammad Jehanzeb Mirza

ICCV 2025

PerLA: Perceptive 3D Language Assistant

Guofeng Mei, Wei Lin, Luigi Riz, Yujiao Wu, Fabio Poiesi, Yiming Wang

CVPR 2025

LiveXiv—A Multi-Modal Live Benchmark Based on Arxiv Papers Content

Nimrod Shabtay, Felipe Maia Polo, Sivan Doveh, Wei Lin, Muhammad Jehanzeb Mirza, Leshem Choshen, Mikhail Yurochkin, Yuekai Sun, Assaf Arbelle, Leonid Karlinsky, Raja Giryes

ICLR 2025

Instructify: Demystifying Metadata to Visual Instruction Tuning Data Conversion

Jacob Hansen, Wei Lin, Junmo Kang, Muhammad Jehanzeb Mirza, Hongyin Luo, Rogerio Feris, Alan Ritter, James Glass, Leonid Karlinsky

Arxiv 2025

GLOV: Guided Large Language Models as Implicit Optimizers for Vision Language Models

Muhammad Jehanzeb Mirza, Mengjie Zhao, Zhuoyuan Mao, Sivan Doveh, Wei Lin, Paul Gavrikov, Michael Dorkenwald, Shiqi Yang, Saurav Jha, Hiromi Wakaki, Yuki Mitsufuji, Horst Possegger, Rogerio Feris, Leonid Karlinsky, James Glass

Arxiv 2024

Comparison Visual Instruction Tuning

Wei Lin, Muhammad Jehanzeb Mirza, Sivan Doveh, Rogerio Feris, Raja Giryes, Sepp Hochreiter, Leonid Karlinsky

In collaboration with the MIT-IBM Watson AI Lab

Arxiv 2024

Conme: Rethinking Evaluation of Compositional Reasoning for Modern VLMs

*Irene Huang, *Wei Lin, *Muhammad Jehanzeb Mirza, Jacob Hansen, Sivan Doveh, Victor Ion Butoi, Roei Herzig, Assaf Arbelle, Hilde Kuehne, Trevor Darrell, Chuang Gan, Aude Oliva, Rogerio Feris, Leonid Karlinsky (*equal contribution)

In collaboration with the MIT-IBM Watson AI Lab

Conference on Neural Information Processing Systems (NeurIPS) 2024

Meta-Prompting for Automating Zero-shot Visual Recognition with LLMs

Muhammad Jehanzeb Mirza, Leonid Karlinsky, Wei Lin, Sivan Doveh, Jakub Micorek, Mateusz Kozinski, Hilde Kuehne, Horst Possegger

In collaboration with the MIT-IBM Watson AI Lab

European Conference on Computer Vision (ECCV) 2024

Towards multimodal in-context learning for vision & language models

Sivan Doveh, Shaked Perek, Muhammad Jehanzeb Mirza, Wei Lin, Amit Alfassy, Assaf Arbelle, Shimon Ullman, Leonid Karlinsky

In collaboration with the MIT-IBM Watson AI Lab

Arxiv 2024

Vision-Language Guidance for LiDAR-based Unsupervised 3D Object Detection

Christian Fruhwirth-Reisinger, Wei Lin, Dusan Malic, Horst Bischof, Horst Possegger

British Machine Vision Conference (BMVC) 2024 Oral & Best Poster Award

Overlooked Aspects in the Evaluation of Out-Of-Distribution Detection Methods

*Bernhard Lehner, *Christian Huber, Bernhard Moser, Claus Hofmann, Wei Lin, Sepp Hochreiter (*equal contribution)

Arxiv 2024

MATch, eXpand and Improve: Unsupervised Finetuning for Zero-Shot Action Recognition with Language Knowledge

Wei Lin, Leonid Karlinsky, Nina Shvetsova, Horst Possegger, Mateusz Kozinski, Rameswar Panda, Rogerio Feris, Hilde Kuehne, Horst Bischof

In collaboration with the MIT-IBM Watson AI Lab

International Conference on Computer Vision (ICCV) 2023

LaFTer: Label-Free Tuning of Zero-shot Classifier using Language and Unlabeled Image Collections

Muhammad Jehanzeb Mirza, Leonid Karlinsky, Wei Lin, Mateusz Kozinski, Horst Possegger, Rogerio Feris, Horst Bischof

Conference on Neural Information Processing Systems (NeurIPS) 2023

MATE: Masked Autoencoders are Online 3D Test-Time Learners

*Muhammad Jehanzeb Mirza, *Inkyu Shin, *Wei Lin, Andreas Schriebl, Kunyang Sun, Jaesung Choe, Horst Possegger, Mateusz Kozinski, In So Kweon, Kun-Jin Yoon, Horst Bischof (*equal contribution)

International Conference on Computer Vision (ICCV) 2023

TAP: Targeted Prompting for Task Adaptive Generation of Textual Training Instances for Visual Classification

Muhammad Jehanzeb Mirza, Leonid Karlinsky, Wei Lin, Horst Possegger, Rogerio Feris, Horst Bischof

Arxiv 2023

Video Test-Time Adaptation for Action Recognition

*Wei Lin, *Muhammad Jehanzeb Mirza, Mateusz Kozinski, Horst Possegger, Hilde Kuehne, Horst Bischof (*equal contribution)

Conference on Computer Vision and Pattern Recognition (CVPR) 2023

ActMAD: Activation Matching to Align Distributions for Test-Time-Training

Muhammad Jehanzeb Mirza, Pol Jané Soneira, Wei Lin, Mateusz Kozinski, Horst Possegger, Horst Bischof

Conference on Computer Vision and Pattern Recognition (CVPR) 2023

CycDA: Unsupervised Cycle Domain Adaptation to Learn from Image to Video

Wei Lin, Anna Kukleva, Kunyang Sun, Horst Possegger, Hilde Kuehne, Horst Bischof

European Conference on Computer Vision (ECCV) 2022

Extended Abstract CycDA: Unsupervised Cycle Domain Adaptation to Learn from Image to Video

Wei Lin, Anna Kukleva, Kunyang Sun, Horst Possegger, Hilde Kuehne, Horst Bischof

ECCV 2022 Workshop of Out Of Distribution Generalization in Computer Vision, 2022

Unsupervised Class-aware 3D Object Detection in LiDAR Point Clouds

Christian Fruhwirth-Reisinger, Wei Lin, Dusan Malic, David Schinagl, Georg Krispel, Horst Possegger, Horst Bischof

Arxiv 2023

AIR-DA: Adversarial Image Reconstruction for Unsupervised Domain Adaptive Object Detection

Kunyang Sun, Wei Lin, Haoqin Shi, Zhengming Zhang, Yongming Huang, Horst Bischof

IEEE Robotics and Automation Letters (RA-L) 2023

TAEC: Unsupervised Action Segmentation with Temporal-Aware Embedding and Clustering

Wei Lin, Anna Kukleva, Horst Possegger, Hilde Kuehne, Horst Bischof

Computer Vision Winter Workshop 2023

Sit Back and Relax: Learning to Drive Incrementally in All Weather Conditions

Stefan Leitner, Muhammad Jehanzeb Mirza, Wei Lin, Jakub Microrek, Marc Masana, Mateusz Kozinski, Horst Possegger, Horst Bischof

Intelligent Vehicle Symposium, 2023

Review service

Conference

ECCV 2022, ISMAR 2023, CVPR 2023, NeurIPS 2023, WACV 2024, CVPR 2024, ECCV 2024, NeurIPS 2024, NeurIPS 2024 Dataset and Benchmark Track, CVPR 2025, ICCV 2025, ICCV 2025 Workshop MMFM4, ICLR 2026, WACV 2026, CVPR 2026, CVPR 2026 Workshop CV4Edu, NeurIPS 2026, ICLR 2027

Journal

TPAMI 2023, TNNLS 2023, IEEE Trans. Multimedia 2023, Pattern Recognition Letters 2024

Activity

International Computer Vision Summer School 2023

Teaching

Deep Learning and Neural Networks I Exercise, Deep Learning and Neural Networks II Exercise, Machine Learning: Supervised Techniques Exercise, Machine Learning: Unsupervised Techniques Exercise

Honors and scholarships

Scholarship for Foreign Students (Technical University of Munich) 06.2016
Scholarship for Foreign Students (Technical University of Munich) 06.2017

Internship

Robert Bosch GmbH, Corporate Research Hildesheim, Germany
Research Internship Oct 2017 - May 2018
Master's Thesis : 3D Human Pose-based Action Recognition

Robert Bosch GmbH, Corporate Research Leonberg, Germany
Research Internship Feb 2017 - Jul 2017
Project: Road surface estimation from monocular video data based on P-Spline regression and 3D reconstruction

Skills

Programming

Proficient in: Python, C++

Languages

English, German, Chinese